

Model

July 29, 2024

```
[1]: import numpy as np  
import tensorflow as tf
```

```
[2]: npz_file = np.load('Audibooks_data_train.npz')  
train_inputs = npz_file['inputs'].astype(float)  
train_targets = npz_file['targets'].astype(int)  
  
npz_file = np.load('Audibooks_data_validation.npz')  
validation_inputs = npz_file['inputs'].astype(float)  
validation_targets = npz_file['targets'].astype(int)  
  
npz_file = np.load('Audibooks_data_test.npz')  
test_inputs = npz_file['inputs'].astype(float)  
test_targets = npz_file['targets'].astype(int)
```

0.0.1 Building and Training Model

```
[3]: input_size = 10  
output_size = 2  
hidden_layer_size = 50  
batch_size = 100  
max_epochs = 100  
  
early_stopping = tf.keras.callbacks.EarlyStopping(patience=2) #Prevent ↴ Overfitting  
  
model = tf.keras.Sequential([  
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),  
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),  
    tf.keras.layers.Dense(output_size, activation='softmax')  
])  
  
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', ↴  
    metrics=['accuracy'])  
  
model.fit(train_inputs, train_targets,  
          batch_size=batch_size,
```

```
epochss=max_epochs,  
callbacks=[early_stopping],  
validation_data=(validation_inputs,validation_targets),  
verbose=2)
```

```
Epoch 1/100  
36/36 - 1s - loss: 0.5790 - accuracy: 0.6806 - val_loss: 0.5090 - val_accuracy:  
0.7651 - 1s/epoch - 29ms/step  
Epoch 2/100  
36/36 - 0s - loss: 0.4611 - accuracy: 0.7714 - val_loss: 0.4434 - val_accuracy:  
0.7830 - 70ms/epoch - 2ms/step  
Epoch 3/100  
36/36 - 0s - loss: 0.4102 - accuracy: 0.7835 - val_loss: 0.4064 - val_accuracy:  
0.7740 - 78ms/epoch - 2ms/step  
Epoch 4/100  
36/36 - 0s - loss: 0.3844 - accuracy: 0.7974 - val_loss: 0.3927 - val_accuracy:  
0.7740 - 91ms/epoch - 3ms/step  
Epoch 5/100  
36/36 - 0s - loss: 0.3697 - accuracy: 0.7944 - val_loss: 0.3827 - val_accuracy:  
0.7919 - 83ms/epoch - 2ms/step  
Epoch 6/100  
36/36 - 0s - loss: 0.3567 - accuracy: 0.8117 - val_loss: 0.3780 - val_accuracy:  
0.7919 - 80ms/epoch - 2ms/step  
Epoch 7/100  
36/36 - 0s - loss: 0.3498 - accuracy: 0.8167 - val_loss: 0.3718 - val_accuracy:  
0.7875 - 79ms/epoch - 2ms/step  
Epoch 8/100  
36/36 - 0s - loss: 0.3433 - accuracy: 0.8131 - val_loss: 0.3638 - val_accuracy:  
0.7942 - 75ms/epoch - 2ms/step  
Epoch 9/100  
36/36 - 0s - loss: 0.3378 - accuracy: 0.8181 - val_loss: 0.3638 - val_accuracy:  
0.8098 - 74ms/epoch - 2ms/step  
Epoch 10/100  
36/36 - 0s - loss: 0.3359 - accuracy: 0.8181 - val_loss: 0.3621 - val_accuracy:  
0.8009 - 74ms/epoch - 2ms/step  
Epoch 11/100  
36/36 - 0s - loss: 0.3332 - accuracy: 0.8170 - val_loss: 0.3583 - val_accuracy:  
0.8031 - 70ms/epoch - 2ms/step  
Epoch 12/100  
36/36 - 0s - loss: 0.3301 - accuracy: 0.8220 - val_loss: 0.3588 - val_accuracy:  
0.8076 - 77ms/epoch - 2ms/step  
Epoch 13/100  
36/36 - 0s - loss: 0.3270 - accuracy: 0.8248 - val_loss: 0.3589 - val_accuracy:  
0.8166 - 78ms/epoch - 2ms/step
```

[3]: <keras.callbacks.History at 0x1f0e3cd2560>

0.0.2 Testing

```
[4]: test_loss, test_accuracy = model.evaluate(test_inputs, test_targets)
```

```
14/14 [=====] - 0s 1ms/step - loss: 0.3275 - accuracy:  
0.8281
```

```
[13]: test_inputs[0], test_inputs[10]
```

```
[13]: (array([ 0.21053387, -0.18888517,  0.34457208,  0.09587919,  2.10780037,  
       -2.60893407,  0.51682305,  0.63990752, -0.20183481, -0.62711546]),  
       array([ 0.21053387, -0.18888517, -0.08811106, -0.25963932, -0.47442823,  
       0.00142655, -0.42647785, -0.41569922, -0.20183481, -0.80255852]))
```

```
[12]: test_targets[0], test_targets[10]
```

```
[12]: (0, 1)
```

```
[8]: y_hat = model.predict(test_inputs)
```

```
14/14 [=====] - 0s 924us/step
```

```
[14]: np.argmax(y_hat[0]), np.argmax(y_hat[10])
```

```
[14]: (0, 1)
```

```
[ ]:
```